



(<https://www.bigdatauniversity.com>)

Classification with Python

In this notebook we try to practice all the classification algorithms that we learned in this course.

We load a dataset using Pandas library, and apply the following algorithms, and find the best one for this specific dataset by accuracy evaluation methods.

Lets first load required libraries:

```
In [1]: import itertools
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.ticker import NullFormatter
import pandas as pd
import numpy as np
import matplotlib.ticker as ticker
from sklearn import preprocessing
%matplotlib inline
```

About dataset

This dataset is about past loans. The **Loan_train.csv** data set includes details of 346 customers whose loan are already paid off or defaulted. It includes following fields:

Field	Description
Loan_status	Whether a loan is paid off on in collection
Principal	Basic principal loan amount at the
Terms	Origination terms which can be weekly (7 days), biweekly, and monthly payoff schedule
Effective_date	When the loan got originated and took effects
Due_date	Since it's one-time payoff schedule, each loan has one single due date
Age	Age of applicant
Education	Education of applicant
Gender	The gender of applicant

Lets download the dataset

```
In [2]: !wget -O loan_train.csv https://s3-api.us-gio.objectstorage.softlayer.net/cf-c
courses-data/CognitiveClass/ML0101ENV3/labs/loan_train.csv

--2019-07-09 22:30:28-- https://s3-api.us-gio.objectstorage.softlayer.net/cf-
courses-data/CognitiveClass/ML0101ENV3/labs/loan_train.csv
Resolving s3-api.us-gio.objectstorage.softlayer.net (s3-api.us-gio.objectstor
age.softlayer.net)... 67.228.254.193
Connecting to s3-api.us-gio.objectstorage.softlayer.net (s3-api.us-gio.object
storage.softlayer.net)|67.228.254.193|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 23101 (23K) [text/csv]
Saving to: 'loan_train.csv'

loan_train.csv      100%[=====>]  22.56K  --.-KB/s    in 0.02s

2019-07-09 22:30:29 (1.04 MB/s) - 'loan_train.csv' saved [23101/23101]
```

Load Data From CSV File

```
In [3]: df = pd.read_csv('loan_train.csv')
df.head()
```

Out[3]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	0	0	PAIDOFF	1000	30	9/8/2016	10/7/2016	45	High School or Below
1	2	2	PAIDOFF	1000	30	9/8/2016	10/7/2016	33	Bechelor
2	3	3	PAIDOFF	1000	15	9/8/2016	9/22/2016	27	college
3	4	4	PAIDOFF	1000	30	9/9/2016	10/8/2016	28	college
4	6	6	PAIDOFF	1000	30	9/9/2016	10/8/2016	29	college

```
In [4]: df.shape
```

Out[4]: (346, 10)

Convert to date time object

```
In [5]: df['due_date'] = pd.to_datetime(df['due_date'])
df['effective_date'] = pd.to_datetime(df['effective_date'])
df.head()
```

Out[5]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	0	0	PAIDOFF	1000	30	2016-09-08	2016-10-07	45	High School or Below
1	2	2	PAIDOFF	1000	30	2016-09-08	2016-10-07	33	Bechalar
2	3	3	PAIDOFF	1000	15	2016-09-08	2016-09-22	27	college
3	4	4	PAIDOFF	1000	30	2016-09-09	2016-10-08	28	college
4	6	6	PAIDOFF	1000	30	2016-09-09	2016-10-08	29	college

Data visualization and pre-processing

Let's see how many of each class is in our data set

```
In [6]: df['loan_status'].value_counts()
```

```
Out[6]: PAIDOFF      260
COLLECTION    86
Name: loan_status, dtype: int64
```

260 people have paid off the loan on time while 86 have gone into collection

Lets plot some columns to underestand data better:

```
In [7]: # notice: installing seaborn might takes a few minutes  
!conda install -c anaconda seaborn -y
```

Collecting package metadata: done

Solving environment: -

The environment is inconsistent, please check the package plan carefully

The following packages are causing the inconsistency:

- anaconda/linux-64::conda-build==3.17.8=py36_0
- anaconda/linux-64::grpcio==1.16.1=py36hf8bcb03_1
- anaconda/linux-64::keras==2.1.5=py36_0
- anaconda/linux-64::libarchive==3.3.3=h5d8350f_5
- anaconda/linux-64::python-libarchive-c==2.8=py36_6
- anaconda/linux-64::tensorboard==1.8.0=py36hf484d3e_0
- anaconda/linux-64::tensorflow==1.8.0=h57681fa_0
- anaconda/linux-64::tensorflow-base==1.8.0=py36h5f64886_0
- defaults/linux-64::anaconda==5.3.1=py37_0
- defaults/linux-64::astropy==3.0.4=py37h14c3975_0
- defaults/linux-64::bkcharts==0.2=py37_0
- defaults/linux-64::blaze==0.11.3=py37_0
- defaults/linux-64::bokeh==0.13.0=py37_0
- defaults/linux-64::bottleneck==1.2.1=py37h035aef0_1
- defaults/linux-64::dask==0.19.1=py37_0
- defaults/linux-64::datashape==0.5.4=py37_1
- defaults/linux-64::mkl-service==1.1.2=py37h90e4bf4_5
- defaults/linux-64::numba==0.39.0=py37h04863e7_0
- defaults/linux-64::numexpr==2.6.8=py37hd89afb7_0
- defaults/linux-64::odo==0.5.1=py37_0
- defaults/linux-64::pytables==3.4.4=py37ha205bf6_0
- defaults/linux-64::pytest-arraydiff==0.2=py37h39e3cac_0
- defaults/linux-64::pytest-astropy==0.4.0=py37_0
- defaults/linux-64::pytest-doctestplus==0.1.3=py37_0
- defaults/linux-64::pywavelets==1.0.0=py37hdd07704_0
- defaults/linux-64::scikit-image==0.14.0=py37hf484d3e_1

done

Package Plan

environment location: /home/jupyterlab/conda

added / updated specs:

- seaborn

The following packages will be downloaded:

package	build		
-----	-----		
ca-certificates-2019.5.15	0	133 KB	anaconda
certifi-2019.6.16	py36_0	154 KB	anaconda
conda-4.7.5	py36_0	3.0 MB	anaconda
conda-package-handling-1.3.10	py36_0	259 KB	anaconda
libtiff-4.0.10	h2733197_2	604 KB	anaconda
python-libarchive-c-2.8	py36_9	22 KB	anaconda
zstd-1.3.7	h0b5b093_0	887 KB	anaconda
-----	-----		
Total:		5.0 MB	

The following NEW packages will be INSTALLED:

```
conda-package-han~ anaconda/linux-64::conda-package-handling-1.3.10-py36_0
```

The following packages will be UPDATED:

```
conda                        4.6.14-py36_0 --> 4.7.5-py36_0
openssl                      conda-forge::openssl-1.1.1b-h14c3975_1 --> anaconda::ope
nssl-1.1.1-h7b6447c_0
python-libarchive~          2.8-py36_6 --> 2.8-py36_9
```

The following packages will be SUPERSEDED by a higher-priority channel:

```
ca-certificates             conda-forge::ca-certificates-2019.6.1~ --> anaconda::ca-
certificates-2019.5.15-0
certifi                      conda-forge --> anaconda
libtiff                      conda-forge::libtiff-4.0.10-h57b8799_~ --> anaconda::lib
tiff-4.0.10-h2733197_2
zstd                         conda-forge::zstd-1.4.0-h3b9ef0a_0 --> anaconda::zst
d-1.3.7-h0b5b093_0
```

Downloading and Extracting Packages

```
python-libarchive-c- | 22 KB      | ##### | 10
0%
libtiff-4.0.10        | 604 KB   | ##### | 10
0%
certifi-2019.6.16     | 154 KB   | ##### | 10
0%
zstd-1.3.7            | 887 KB   | ##### | 10
0%
conda-4.7.5           | 3.0 MB   | ##### | 10
0%
ca-certificates-2019 | 133 KB   | ##### | 10
0%
conda-package-handli | 259 KB   | ##### | 10
0%
```

Preparing transaction: done


Verifying transaction: done

Executing transaction: done

```
In [8]: import seaborn as sns


bins = np.linspace(df.Principal.min(), df.Principal.max(), 10)
g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
g.map(plt.hist, 'Principal', bins=bins, ec="k")

g.axes[-1].legend()
plt.show()
```



```
In [9]: bins = np.linspace(df.age.min(), df.age.max(), 10)
g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
g.map(plt.hist, 'age', bins=bins, ec="k")


g.axes[-1].legend()
plt.show()
```



Pre-processing: Feature selection/extraction

Lets look at the day of the week people get the loan

```
In [10]: df['dayofweek'] = df['effective_date'].dt.dayofweek
bins = np.linspace(df.dayofweek.min(), df.dayofweek.max(), 10)
g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
g.map(plt.hist, 'dayofweek', bins=bins, ec="k")
g.axes[-1].legend()
plt.show()
```



We see that people who get the loan at the end of the week dont pay it off, so lets use Feature binarization to set a threshold values less then day 4

```
In [11]: df['weekend'] = df['dayofweek'].apply(lambda x: 1 if (x>3) else 0)
df.head()
```

Out[11]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	0	0	PAIDOFF	1000	30	2016-09-08	2016-10-07	45	High School or Below
1	2	2	PAIDOFF	1000	30	2016-09-08	2016-10-07	33	Bechalar
2	3	3	PAIDOFF	1000	15	2016-09-08	2016-09-22	27	college
3	4	4	PAIDOFF	1000	30	2016-09-09	2016-10-08	28	college
4	6	6	PAIDOFF	1000	30	2016-09-09	2016-10-08	29	college

Convert Categorical features to numerical values

Lets look at gender:


```
In [12]: df.groupby(['Gender'])['loan_status'].value_counts(normalize=True)
```

```
Out[12]: Gender  loan_status
female  PAIDOFF      0.865385
        COLLECTION  0.134615
male    PAIDOFF      0.731293
        COLLECTION  0.268707
Name: loan_status, dtype: float64
```

86 % of female pay there loans while only 73 % of males pay there loan

Lets convert male to 0 and female to 1:

```
In [13]: df['Gender'].replace(to_replace=['male','female'], value=[0,1],inplace=True)
df.head()
```

```
Out[13]:
```

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	0	0	PAIDOFF	1000	30	2016-09-08	2016-10-07	45	High School or Below
1	2	2	PAIDOFF	1000	30	2016-09-08	2016-10-07	33	Bechalar
2	3	3	PAIDOFF	1000	15	2016-09-08	2016-09-22	27	college
3	4	4	PAIDOFF	1000	30	2016-09-09	2016-10-08	28	college
4	6	6	PAIDOFF	1000	30	2016-09-09	2016-10-08	29	college

One Hot Encoding

How about education?

```
In [14]: df.groupby(['education'])['loan_status'].value_counts(normalize=True)
```

```
Out[14]: education  loan_status
Bechalar          PAIDOFF      0.750000
                  COLLECTION  0.250000
High School or Below PAIDOFF      0.741722
                  COLLECTION  0.258278
Master or Above    COLLECTION  0.500000
                  PAIDOFF      0.500000
college           PAIDOFF      0.765101
                  COLLECTION  0.234899
Name: loan_status, dtype: float64
```

Feature before One Hot Encoding

```
In [15]: df[['Principal', 'terms', 'age', 'Gender', 'education']].head()
```

Out[15]:

	Principal	terms	age	Gender	education
0	1000	30	45	0	High School or Below
1	1000	30	33	1	Bechalor
2	1000	15	27	0	college
3	1000	30	28	1	college
4	1000	30	29	0	college

Use one hot encoding technique to conver categorical variables to binary variables and append them to the feature Data Frame

```
In [16]: Feature = df[['Principal', 'terms', 'age', 'Gender', 'weekend']]
Feature = pd.concat([Feature, pd.get_dummies(df['education'])], axis=1)
Feature.drop(['Master or Above'], axis = 1, inplace=True)
Feature.head()
```

Out[16]:

	Principal	terms	age	Gender	weekend	Bechalor	High School or Below	college
0	1000	30	45	0	0	0	1	0
1	1000	30	33	1	0	1	0	0
2	1000	15	27	0	0	0	0	1
3	1000	30	28	1	1	0	0	1
4	1000	30	29	0	1	0	0	1

Feature selection

Lets definid feature sets, X:

```
In [17]: X_train=Feature
X = Feature
X[0:5]
```

Out[17]:

	Principal	terms	age	Gender	weekend	Bechalor	High School or Below	college
0	1000	30	45	0	0	0	1	0
1	1000	30	33	1	0	1	0	0
2	1000	15	27	0	0	0	0	1
3	1000	30	28	1	1	0	0	1
4	1000	30	29	0	1	0	0	1

What are our lables?

```
In [18]: y_train = df['loan_status'].values
y = df['loan_status'].values
y[0:5]
```

Out[18]: array(['PAIDOFF', 'PAIDOFF', 'PAIDOFF', 'PAIDOFF', 'PAIDOFF'],
dtype=object)

Normalize Data

Data Standardization give data zero mean and unit variance (technically should be done after train test split)

```
In [19]: X_train= preprocessing.StandardScaler().fit(X).transform(X)
X_train[0:5]
```

```
/home/jupyterlab/conda/lib/python3.6/site-packages/sklearn/preprocessing/data
a.py:625: DataConversionWarning: Data with input dtype uint8, int64 were all
converted to float64 by StandardScaler.
    return self.partial_fit(X, y)
/home/jupyterlab/conda/lib/python3.6/site-packages/ipykernel_launcher.py:1: D
ataConversionWarning: Data with input dtype uint8, int64 were all converted t
o float64 by StandardScaler.
    """Entry point for launching an IPython kernel.
```

Out[19]: array([[0.51578458, 0.92071769, 2.33152555, -0.42056004, -1.20577805,
-0.38170062, 1.13639374, -0.86968108],
[0.51578458, 0.92071769, 0.34170148, 2.37778177, -1.20577805,
2.61985426, -0.87997669, -0.86968108],
[0.51578458, -0.95911111, -0.65321055, -0.42056004, -1.20577805,
-0.38170062, -0.87997669, 1.14984679],
[0.51578458, 0.92071769, -0.48739188, 2.37778177, 0.82934003,
-0.38170062, -0.87997669, 1.14984679],
[0.51578458, 0.92071769, -0.3215732 , -0.42056004, 0.82934003,
-0.38170062, -0.87997669, 1.14984679]])

Classification

In []:

Now, it is your turn, use the training set to build an accurate model. Then use the test set to report the accuracy of the model You should use the following algorithm:

- K Nearest Neighbor(KNN)
- Decision Tree
- Support Vector Machine
- Logistic Regression

Notice:

- You can go above and change the pre-processing, feature selection, feature-extraction, and so on, to make a better model.
- You should use either scikit-learn, Scipy or Numpy libraries for developing the classification algorithms.
- You should include the code of the algorithm in the following cells.

K Nearest Neighbor(KNN)

Notice: You should find the best k to build the model with the best accuracy.

warning: You should not use the **loan_test.csv** for finding the best k, however, you can split your train_loan.csv into train and test to find the best k.

```
In [20]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [21]: k = 4
#Train Model and Predict
neigh = KNeighborsClassifier(n_neighbors = k).fit(X_train,y_train)
neigh
```

```
Out[21]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=4, p=2,
                             weights='uniform')
```

In []:

Decision Tree

```
In [22]: from sklearn.tree import DecisionTreeClassifier
```

```
In [23]: loanTree = DecisionTreeClassifier(criterion="entropy", max_depth = 4)
loanTree # it shows the default parameters
```

```
Out[23]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=4,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best')
```

```
In [24]: loanTree.fit(X_train,y_train)
```

```
Out[24]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=4,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best')
```

Support Vector Machine

```
In [25]: from sklearn import svm
clf = svm.SVC(kernel='rbf')
clf.fit(X_train, y_train)
```

```
Out[25]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
kernel='rbf', max_iter=-1, probability=False, random_state=None,
shrinking=True, tol=0.001, verbose=False)
```

```
In [ ]:
```

```
In [ ]:
```

Logistic Regression

```
In [26]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
LR = LogisticRegression(C=0.01, solver='liblinear').fit(X_train,y_train)
LR
```

```
Out[26]: LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='warn',
n_jobs=None, penalty='l2', random_state=None, solver='liblinear',
tol=0.0001, verbose=0, warm_start=False)
```

```
In [ ]:
```

In []:

Model Evaluation using Test set

```
In [27]: from sklearn.metrics import jaccard_similarity_score
from sklearn.metrics import f1_score
from sklearn.metrics import log_loss
```

First, download and load the test set:

```
In [28]: !wget -O loan_test.csv https://s3-api.us-gio.objectstorage.softlayer.net/cf-co
urses-data/CognitiveClass/ML0101ENV3/labs/loan_test.csv
```

```
--2019-07-09 22:46:39-- https://s3-api.us-gio.objectstorage.softlayer.net/cf-
-courses-data/CognitiveClass/ML0101ENV3/labs/loan_test.csv
Resolving s3-api.us-gio.objectstorage.softlayer.net (s3-api.us-gio.objectstor
age.softlayer.net)... 67.228.254.193
Connecting to s3-api.us-gio.objectstorage.softlayer.net (s3-api.us-gio.object
storage.softlayer.net)|67.228.254.193|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 3642 (3.6K) [text/csv]
Saving to: 'loan_test.csv'
```

```
loan_test.csv      100%[=====>]    3.56K  --.-KB/s    in 0s
```

```
2019-07-09 22:46:39 (57.1 MB/s) - 'loan_test.csv' saved [3642/3642]
```

Load Test set for evaluation

```
In [36]: test_df = pd.read_csv('loan_test.csv')
test_df.head()
```

Out[36]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	1	1	PAIDOFF	1000	30	9/8/2016	10/7/2016	50	Bechelor
1	5	5	PAIDOFF	300	7	9/9/2016	9/15/2016	35	Master or Above
2	21	21	PAIDOFF	1000	30	9/10/2016	10/9/2016	43	High School or Below
3	24	24	PAIDOFF	1000	30	9/10/2016	10/9/2016	26	college
4	35	35	PAIDOFF	800	15	9/11/2016	9/25/2016	29	Bechelor

Subset the same variables, convert to dummy values and Normalize Data

```
In [37]: test_df['due_date'] = pd.to_datetime(test_df['due_date'])
test_df['effective_date'] = pd.to_datetime(test_df['effective_date'])
test_df.head()
```

Out[37]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	1	1	PAIDOFF	1000	30	2016-09-08	2016-10-07	50	Bechalar
1	5	5	PAIDOFF	300	7	2016-09-09	2016-09-15	35	Master or Above
2	21	21	PAIDOFF	1000	30	2016-09-10	2016-10-09	43	High School or Below
3	24	24	PAIDOFF	1000	30	2016-09-10	2016-10-09	26	college
4	35	35	PAIDOFF	800	15	2016-09-11	2016-09-25	29	Bechalar

```
In [38]: test_df['dayofweek'] = test_df['effective_date'].dt.dayofweek
test_df['weekend'] = test_df['dayofweek'].apply(lambda x: 1 if (x>3) else 0)
test_df['Gender'].replace(to_replace=['male','female'], value=[0,1],inplace=True)
test_df.head()
```

Out[38]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	1	1	PAIDOFF	1000	30	2016-09-08	2016-10-07	50	Bechalar
1	5	5	PAIDOFF	300	7	2016-09-09	2016-09-15	35	Master or Above
2	21	21	PAIDOFF	1000	30	2016-09-10	2016-10-09	43	High School or Below
3	24	24	PAIDOFF	1000	30	2016-09-10	2016-10-09	26	college
4	35	35	PAIDOFF	800	15	2016-09-11	2016-09-25	29	Bechalar

```
In [39]: Feat_t= test_df[['Principal','terms','age','Gender','weekend']] # 'weekend'
Feat_t = pd.concat([Feat_t,pd.get_dummies(test_df['education'])], axis=1)
Feat_t.drop(['Master or Above'], axis = 1,inplace=True)
Feat_t.head()
```

Out[39]:

	Principal	terms	age	Gender	weekend	Bechalor	High School or Below	college
0	1000	30	50	1	0	1	0	0
1	300	7	35	0	1	0	0	0
2	1000	30	43	1	1	0	1	0
3	1000	30	26	0	1	0	0	1
4	800	15	29	0	1	1	0	0

```
In [40]: X_test=Feat_t
X_test= preprocessing.StandardScaler().fit(X_test).transform(X_test)
X_test[0:5]
```

/home/jupyterlab/conda/lib/python3.6/site-packages/sklearn/preprocessing/data.py:625: DataConversionWarning: Data with input dtype uint8, int64 were all converted to float64 by StandardScaler.
return self.partial_fit(X, y)
/home/jupyterlab/conda/lib/python3.6/site-packages/ipykernel_launcher.py:2: DataConversionWarning: Data with input dtype uint8, int64 were all converted to float64 by StandardScaler.

```
Out[40]: array([[ 0.49362588,  0.92844966,  3.05981865,  1.97714211, -1.30384048,
  2.39791576, -0.79772404, -0.86135677],
 [-3.56269116, -1.70427745,  0.53336288, -0.50578054,  0.76696499,
 -0.41702883, -0.79772404, -0.86135677],
 [ 0.49362588,  0.92844966,  1.88080596,  1.97714211,  0.76696499,
 -0.41702883,  1.25356634, -0.86135677],
 [ 0.49362588,  0.92844966, -0.98251057, -0.50578054,  0.76696499,
 -0.41702883, -0.79772404,  1.16095912],
 [-0.66532184, -0.78854628, -0.47721942, -0.50578054,  0.76696499,
  2.39791576, -0.79772404, -0.86135677]])
```

```
In [41]: y_test = test_df['loan_status'].values
y_test[0:5]
```

```
Out[41]: array(['PAIDOFF', 'PAIDOFF', 'PAIDOFF', 'PAIDOFF', 'PAIDOFF'],
              dtype=object)
```

KNN

```
In [42]: yhat = neigh.predict(X_test)
yhat[0:5]
```

```
Out[42]: array(['PAIDOFF', 'COLLECTION', 'COLLECTION', 'COLLECTION', 'COLLECTION'],
              dtype=object)
```



```
In [45]: print("Jaccar: ",jaccard_similarity_score(y_test, yhat))
print("F1 Score: ",f1_score(y_test, yhat, average='weighted'))
```

Jaccar: 0.6296296296296297
F1 Score: 0.6430311890838205

Decision Tree

```
In [47]: print("Jaccar: ",jaccard_similarity_score(y_test,loanTree.predict(X_test)))
print("F1 Score: ",f1_score(y_test,loanTree.predict(X_test), average='weighted'))
```

Jaccar: 0.7777777777777778
F1 Score: 0.7283950617283951

SVM

```
In [48]: print("Jaccar: ",jaccard_similarity_score(y_test,clf.predict(X_test)))
print("F1 Score: ",f1_score(y_test,clf.predict(X_test), average='weighted'))
```

Jaccar: 0.7222222222222222
F1 Score: 0.6212664277180406

In []:

Logistic Regression

```
In [51]: yhat_prob = LR.predict_proba(X_test)
yhat_prob[0:5]
```

```
Out[51]: array([[0.25256814, 0.74743186],
[0.40233132, 0.59766868],
[0.42774804, 0.57225196],
[0.47276992, 0.52723008],
[0.44726818, 0.55273182]])
```

```
In [52]: print("Jaccar: ",jaccard_similarity_score(y_test,LR.predict(X_test)))
print("F1 Score: ",f1_score(y_test,LR.predict(X_test), average='weighted'))
print("Log Losss: ",log_loss(y_test, yhat_prob))
```

Jaccar: 0.7407407407407407
F1 Score: 0.6304176516942475
Log Losss: 0.5566084946309207

/home/jupyterlab/conda/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: F-score is ill-defined and being set to 0.0 in labels with no predicted samples.
'precision', 'predicted', average, warn_for)

Report

You should be able to report the accuracy of the built model using different evaluation metrics:

Algorithm	Jaccard	F1-score	LogLoss
KNN	0.62962	0.64303	NA
Decision Tree	0.77777	0.72839	NA
SVM	0.72727	0.62126	NA
LogisticRegression	0.74074	0.63041	0.55660

Want to learn more?

IBM SPSS Modeler is a comprehensive analytics platform that has many machine learning algorithms. It has been designed to bring predictive intelligence to decisions made by individuals, by groups, by systems – by your enterprise as a whole. A free trial is available through this course, available here: [SPSS Modeler \(http://cocl.us/ML0101EN-SPSSModeler\)](http://cocl.us/ML0101EN-SPSSModeler)

Also, you can use Watson Studio to run these notebooks faster with bigger datasets. Watson Studio is IBM's leading cloud solution for data scientists, built by data scientists. With Jupyter notebooks, RStudio, Apache Spark and popular libraries pre-packaged in the cloud, Watson Studio enables data scientists to collaborate on their projects without having to install anything. Join the fast-growing community of Watson Studio users today with a free account at [Watson Studio \(https://cocl.us/ML0101EN_DSX\)](https://cocl.us/ML0101EN_DSX)

Thanks for completing this lesson!

Author: [Saeed Aghabozorgi \(https://ca.linkedin.com/in/saeedaghabozorgi\)](https://ca.linkedin.com/in/saeedaghabozorgi)

[Saeed Aghabozorgi \(https://ca.linkedin.com/in/saeedaghabozorgi\)](https://ca.linkedin.com/in/saeedaghabozorgi), PhD is a Data Scientist in IBM with a track record of developing enterprise level applications that substantially increases clients' ability to turn data into actionable knowledge. He is a researcher in data mining field and expert in developing advanced analytic methods like machine learning and statistical modelling on large datasets.

Copyright © 2018 [Cognitive Class \(https://cocl.us/DX0108EN_CC\)](https://cocl.us/DX0108EN_CC). This notebook and its source code are released under the terms of the [MIT License \(https://bigdatauniversity.com/mit-license/\)](https://bigdatauniversity.com/mit-license/).